

Scripts Execution

Explanation of the solution to the batch layer problem

Step by step process followed for the completion of tasks till **task 4**. The steps are documented here.

Problem Statement: Credit card fraud is defined as a form of identity theft in which an individual uses someone else's credit card information to make purchase or to withdraw funds from the account. It can be Fraudulent transactions or customer information.

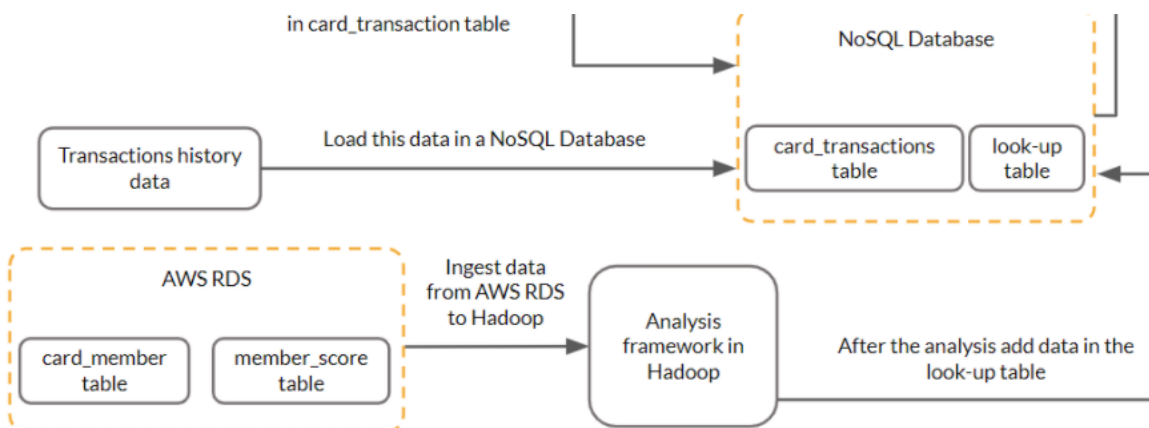
We have **following tasks** to be performed as per **batch layer problem** (till task4)

Task1: Load the transactions history data (card_transactions.csv) in a NoSQL database

Task2: Ingest the relevant data from AWS RDS to Hadoop

Task3: Create a lookup-table with columns specified

Task4: Load the data in lookup table



Task1: Load the transactions history data (card_transactions.csv) in a NoSQL database

Steps:

1. Downloaded the card_transactions.csv from the resources
2. Copied to ec2-user local folder
3. Copied from ec2-user to hdfs (/user/root/capstone_project) using `hadoop fs -put <src> <dest>`
4. Using Pyspark read this csv file and created a dataframe

5. Created hbase table for card_transactions
6. Loaded the data from card_transactions.csv to card_transactions table in hbase
7. Used 'list' command in hbase to see the table created
8. Used scan "card_transactions" to check on the data
9. Used count "card_transactions" to check on count of data.

```
root@ip-10-0-0-149:~
login as: ec2-user
Authenticating with public key "imported-openssh-key"
Last login: Sat Oct 16 08:11:14 2021 from 106.202.196.118
[ec2-user@ip-10-0-0-149 ~]$ ls
card_transactions.csv
[ec2-user@ip-10-0-0-149 ~]$ sudo -i
[root@ip-10-0-0-149 ~]# hadoop fs -put /home/ec2-user/card_transactions.csv /user/root/capstone_project
[root@ip-10-0-0-149 ~]#
```

```
root@ip-10-0-0-138:~
[ec2-user@ip-10-0-0-138 ~]# hadoop fs -ls /user/root/capstone_proj
Found 3 items
drwxr-xr-x - root supergroup 0 2021-10-10 06:40 /user/root/capstone_proj/card_member
-rw-r--r-- 3 root supergroup 4829520 2021-10-10 15:13 /user/root/capstone_proj/card_transactions.csv
drwxr-xr-x - root supergroup 0 2021-10-10 06:38 /user/root/capstone_proj/member_score
[root@ip-10-0-0-138 ~]#
```

```
transasction = StructType([StructField('card_id', StringType(), False),
                               StructField('member_id', StringType(), False),
                               StructField('amount', IntegerType(), False),
                               StructField('postcode', StringType(), False),
                               StructField('pos_id', StringType(), False),
                               StructField('transaction_dt', StringType(), False),
                               StructField('status', StringType(), False),
                               ])
```

Reading Past Transactions data (source as csv)

```
trans_df = spark.read.csv("hdfs:/user/root/capstone_proj/card_transactions.csv", header = True, schema = transasction)
```

```
# create table for card_transactions.csv file.
create_table('card_transactions', {'info' : dict(max_versions=5) })
```

```
creating table card_transactions
fetching all table
all tables fetched
table created
```

```
#Batch insert data of card_transactions.csv file.
batch_insert_csvdata('card_transactions.csv','card_transactions')
```

```
starting batch insert of events
batch insert done
```

```
hbase(main):002:0> list
TABLE
card_transactions
lookup_table
2 row(s) in 0.0080 seconds

=> ["card_transactions", "lookup_table"]
hbase(main):003:0>
```

root@ip-10-0-0-149:~

```
6591175617713393 column=info:transaction_date, timestamp=1634375007372, value=2018-01-31 13:10:37
6592184145413632 column=info:UCL, timestamp=1634375006787, value=13734342.65
6592184145413632 column=info:card_id, timestamp=1634375006787, value=6592184145413632
6592184145413632 column=info:postcode, timestamp=1634375006787, value=53186
6592184145413632 column=info:score, timestamp=1634375006787, value=456
6592184145413632 column=info:transaction_date, timestamp=1634375006787, value=2018-01-28 00:54:30
6594248319343442 column=info:UCL, timestamp=1634375006872, value=15065362.77
6594248319343442 column=info:card_id, timestamp=1634375006872, value=6594248319343442
6594248319343442 column=info:postcode, timestamp=1634375006872, value=24927
6594248319343442 column=info:score, timestamp=1634375006872, value=350
6594248319343442 column=info:transaction_date, timestamp=1634375006872, value=2018-01-31 23:42:38
6595638658736751 column=info:UCL, timestamp=1634375007621, value=14005069.97
6595638658736751 column=info:card_id, timestamp=1634375007621, value=6595638658736751
6595638658736751 column=info:postcode, timestamp=1634375007621, value=68328
6595638658736751 column=info:score, timestamp=1634375007621, value=310
6595638658736751 column=info:transaction_date, timestamp=1634375007621, value=2018-01-30 10:50:34
6595814135833988 column=info:UCL, timestamp=1634375007288, value=14332708.84
6595814135833988 column=info:card_id, timestamp=1634375007288, value=6595814135833988
6595814135833988 column=info:postcode, timestamp=1634375007288, value=22508
6595814135833988 column=info:score, timestamp=1634375007288, value=210
6595814135833988 column=info:transaction_date, timestamp=1634375007288, value=2018-01-30 02:03:54
6595928469079750 column=info:UCL, timestamp=1634375008323, value=11824730.01
6595928469079750 column=info:card_id, timestamp=1634375008323, value=6595928469079750
6595928469079750 column=info:postcode, timestamp=1634375008323, value=98349
6595928469079750 column=info:score, timestamp=1634375008323, value=412
6595928469079750 column=info:transaction_date, timestamp=1634375008323, value=2018-01-24 12:38:22
6597703848279563 column=info:UCL, timestamp=1634375007681, value=15250624.49
6597703848279563 column=info:card_id, timestamp=1634375007681, value=6597703848279563
6597703848279563 column=info:postcode, timestamp=1634375007681, value=95699
6597703848279563 column=info:score, timestamp=1634375007681, value=218
6597703848279563 column=info:transaction_date, timestamp=1634375007681, value=2018-01-27 10:51:49
6598830758632447 column=info:UCL, timestamp=1634375007878, value=12685782.48
6598830758632447 column=info:card_id, timestamp=1634375007878, value=6598830758632447
6598830758632447 column=info:postcode, timestamp=1634375007878, value=19421
6598830758632447 column=info:score, timestamp=1634375007878, value=293
6598830758632447 column=info:transaction_date, timestamp=1634375007878, value=2018-01-30 00:18:34
6599900931314251 column=info:UCL, timestamp=1634375008288, value=12487392.07
6599900931314251 column=info:card_id, timestamp=1634375008288, value=6599900931314251
6599900931314251 column=info:postcode, timestamp=1634375008288, value=97423
6599900931314251 column=info:score, timestamp=1634375008288, value=297
6599900931314251 column=info:transaction_date, timestamp=1634375008288, value=2018-01-31 11:25:16
999 row(s) in 0.4250 seconds
hbase(main):009:0>
```

Task2: Ingest the relevant data from AWS RDS to Hadoop

Steps:

1. Checked the details for ingesting data from resources.
2. Used sqoop import for importing data to hadoop

Table 1:member_score

```
sqoop import \  
--connect jdbc:mysql://upgradawsrds1.cyaieic9bmnf.us-  
east1.rds.amazonaws.com/cred_financials_data \  
--table member_score \  
--username upgraduser --password upgraduser \  
--target-dir /user/root/capstone_project/member_score \  
-m 1
```

Table 2: card_member

```
sqoop import \  
--connect jdbc:mysql://upgradawsrds1.cyaieic9bmnf.us-  
east1.rds.amazonaws.com/cred_financials_data \  
--table card_member \  
--username upgraduser --password upgraduser \  
--target-dir /user/root/capstone_project/card_member \  
-m 1
```

```
[root@ip-10-0-0-149 ~]# hadoop fs -ls /user/root/capstone_project  
Found 2 items  
drwxr-xr-x - root supergroup 0 2021-10-16 07:29 /user/root/capstone_project/card_member  
drwxr-xr-x - root supergroup 0 2021-10-16 07:23 /user/root/capstone_project/member_score
```

Task3: Create a lookup-table with columns specified

Steps:

1. Once the card_member and member_score data is available in task2 is available in Pyspark as dataframes
2. Used this as a source and merged databases for relevant columns
3. Calculated extra columns required and created lookup table in hbase

```
cardschema = StructType([StructField('card_id', StringType(),False),
                             StructField('member_id', StringType(),False),
                             StructField('member_joining_dt', StringType(),False),
                             StructField('card_purchase_dt', StringType(),False),
                             StructField('country', StringType(),False),
                             StructField('city', StringType(),False),
                             ])
```

#read the data

```
card_df = spark.read.csv("hdfs:/user/root/capstone_project/card_member", header = False, schema = cardschema)
```

card_id	member_id	member_joining_dt	card_purchase_dt	country	city
340028465709212	009250698176266	2012-02-08 06:04:...	05/13	United States	Barberton
340054675199675	835873341185231	2017-03-10 09:24:...	03/17	United States	Fort Dodge
340082915339645	512969555857346	2014-02-15 06:30:...	07/14	United States	Graham
340134186926007	887711945571282	2012-02-05 01:21:...	02/13	United States	Dix Hills
340265728490548	680324265406190	2014-03-29 07:49:...	11/14	United States	Rancho Cucamonga
340268219434811	929799084911715	2012-07-08 02:46:...	08/12	United States	San Francisco
340379737226464	089615510858348	2010-03-10 00:06:...	09/10	United States	Clinton
340383645652108	181180599313885	2012-02-24 05:32:...	10/16	United States	West New York
340803866934451	417664728506297	2015-05-21 04:30:...	08/17	United States	Beaverton
340889618969736	459292914761635	2013-04-23 08:40:...	11/15	United States	West Palm Beach
340924125838453	188119365574843	2011-04-12 04:28:...	12/13	United States	Scottsbluff
341005627432127	872138964937565	2013-09-08 03:16:...	02/17	United States	Chillum
341029651579925	974087224071871	2011-01-14 00:20:...	08/12	United States	Valley Station
341311317050937	561687420200207	2014-03-18 06:23:...	02/15	United States	Vincennes
341344252914274	695906467918552	2012-03-02 03:21:...	03/13	United States	Columbine
341363858179050	009190444424572	2012-02-19 05:16:...	04/14	United States	Cheektowaga
341519629171378	533670008048847	2013-05-13 07:59:...	01/15	United States	Centennial
341641153427489	230523184584316	2013-03-25 08:51:...	11/15	United States	Colchester
341719092861087	304847505155781	2015-12-06 08:06:...	11/17	United States	Vernon Hills
341722035429601	979218131207765	2015-12-22 10:46:...	01/17	United States	Elk Grove Village

only showing top 20 rows

```
memberschema = StructType([StructField('member_id', StringType(),False),
                             StructField('score', IntegerType(),False),
                             ])
```

#read the data

```
mem_df = spark.read.csv("hdfs:/user/root/capstone_project/member_score", header = False, schema = memberschema)
```

```
mem_df.show()
```

member_id	score
000037495066290	339
000117826301530	289
001147922084344	393
001314074991813	225
001739553947511	642
003761426295463	413
004494068832701	217
006836124210484	504
006991872634058	697
007955566230397	372
008732267588672	213
008765307152821	399
009136568025042	308
009190444424572	559
009250698176266	233
009873334520465	298
011716573646690	249
011877954983420	497
012390918683920	407
012731668664932	612

only showing top 20 rows

Join the card_member and member_score tables to extract credit score of each member

```
score = mem_df.join(card_df, mem_df.mem_id == card_df.member_id, how='LEFT')
```

```
score.count()
```

```
999
```

```
score.printSchema()
```

```
root
|-- mem_id: string (nullable = true)
|-- score: integer (nullable = true)
|-- card_id: string (nullable = true)
|-- member_id: string (nullable = true)
|-- member_joining_dt: string (nullable = true)
|-- card_purchase_dt: string (nullable = true)
|-- country: string (nullable = true)
|-- city: string (nullable = true)
```

Extract required columns from the joined dataframe

```
score.show()
```

```
+-----+-----+-----+
|      mem_id|score|      cardid|
+-----+-----+-----+
|000037495066290|  339| 348702330256514|
|000117826301530|  289| 5189563368503974|
|001147922084344|  393| 5407073344486464|
|001314074991813|  225| 378303738095292|
|001739553947511|  642| 348413196172048|
|003761426295463|  413| 348536585266345|
|004494068832701|  217| 5515987071565183|
|006836124210484|  504| 5400251558458125|
|006991872634058|  697| 4573337022888445|
|007955566230397|  372| 4708912758619517|
|008732267588672|  213| 5342400571435088|
|008765307152821|  399| 4237648081700588|
|009136568025042|  308| 371814781663843|
|009190444424572|  559| 341363858179050|
|009250698176266|  233| 340028465709212|
|009873334520465|  298| 5495445301620991|
|011716573646690|  249| 4795844193055110|
|011877954983420|  497| 5164771396791995|
|012390918683920|  407| 5423921058459194|
|012731668664932|  612| 5379610024035907|
+-----+-----+-----+
```

only showing top 20 rows

Join both Transaction history and score Dataframe which is a merged and extracted data frame from both RDS tables

```
hist = trans_df.join(score, trans_df.member_id == score.mem_id,how='outer') #outer join on member_id between score ar
hist.count() #53210
53210
hist.printSchema()
root
|-- card_id: string (nullable = true)
|-- member_id: string (nullable = true)
|-- amount: integer (nullable = true)
|-- postcode: string (nullable = true)
|-- pos_id: string (nullable = true)
|-- transaction_dt: string (nullable = true)
|-- status: string (nullable = true)
|-- mem_id: string (nullable = true)
|-- score: integer (nullable = true)
|-- cardid: string (nullable = true)
```



```
hist = hist.select('card_id', 'amount', 'postcode', 'pos_id', 'transaction_dt', 'status', 'score')
```

```
hist.show()
```

card_id	amount	postcode	pos_id	transaction_dt	status	score
340379737226464	6126197	46933	167473544283898	01-05-2016 08:10:50	GENUINE	229
340379737226464	7949232	61840	664980919335952	01-10-2016 10:38:52	GENUINE	229
340379737226464	943839	91743	633038040069180	02-08-2016 00:31:25	GENUINE	229
340379737226464	3764114	91743	633038040069180	02-08-2016 21:35:27	GENUINE	229
340379737226464	6221251	98384	064948657945290	02-10-2016 14:44:14	GENUINE	229
340379737226464	2868312	26032	856772774421259	02-12-2016 21:55:43	GENUINE	229
340379737226464	4418586	20129	390339673634463	02-12-2017 17:05:51	GENUINE	229
340379737226464	7439113	91763	315067016872305	03-04-2017 11:43:59	GENUINE	229
340379737226464	8217180	16063	208378790148728	03-05-2017 16:47:43	GENUINE	229
340379737226464	8505852	64070	695556848392133	03-06-2017 03:07:27	GENUINE	229
340379737226464	8535431	29817	683602833507395	04-08-2016 20:59:31	GENUINE	229
340379737226464	6317993	28425	258522244165233	05-05-2017 00:23:45	GENUINE	229
340379737226464	3256860	16845	933410474855991	05-10-2017 15:09:09	GENUINE	229
340379737226464	1423779	97640	789378980336517	06-02-2017 02:10:00	GENUINE	229
340379737226464	3783517	70552	963177679534627	06-12-2016 03:10:30	GENUINE	229
340379737226464	3300714	75750	072728631441941	07-01-2017 05:52:58	GENUINE	229
340379737226464	5706163	50455	915439934619047	07-01-2018 22:07:07	GENUINE	229
340379737226464	7445128	50455	915439934619047	07-01-2018 23:52:27	GENUINE	229
340379737226464	140120	18915	691571327905821	07-02-2017 20:18:04	GENUINE	229
340379737226464	7720484	48423	548702836055067	07-03-2016 14:59:35	GENUINE	229

only showing top 20 rows

To calculate the latest transaction date of that card:

- group the merged dataset on card_id
- aggregate to max of transaction date.
- Alias the aggregated date as transaction_date

card_id	transaction_date
340379737226464	2018-01-27 00:19:47
377201318164757	2017-11-28 16:32:22
348962542187595	2018-01-29 17:17:14
4389973676463558	2018-01-26 13:47:46
5403923427969691	2018-01-22 23:46:19
345406224887566	2017-12-25 04:03:58
6562510549485881	2018-01-17 08:35:27
5508842242491554	2018-01-31 14:55:58
4407230633003235	2018-01-27 07:21:08
379321864695232	2018-01-03 00:29:37
340028465709212	2018-01-02 03:25:35
349143706735646	2018-01-29 22:33:14
4126356979547079	2018-01-24 16:09:03
5543219113990484	2018-01-13 18:34:00
5464688416792307	2018-01-26 19:03:47
6011273561157733	2018-02-01 01:27:58
4484950467600170	2018-01-10 08:03:13
4818950814628962	2018-01-31 00:53:15
5573293264792992	2018-01-31 14:55:57
6011985140563103	2018-01-30 02:03:54

only showing top 20 rows

Join previous last step data frame (score) with look_up_table dataset created above. This step frames all required cols for look_up_table except the UCL.

```
lookup_table = lookup_table.join(score, lookup_table.card_id == score.cardid, how='INNER')
```

```
lookup_table.count() #check the count (999)
```

999

```
lookup_table.show()
```

card_id	transaction_date	mem_id	score	cardid
340379737226464	2018-01-27 00:19:47	089615510858348	229	340379737226464
345406224887566	2017-12-25 04:03:58	296206661780881	349	345406224887566
348962542187595	2018-01-29 17:17:14	366246487993992	522	348962542187595
377201318164757	2017-11-28 16:32:22	924475891017022	432	377201318164757
379321864695232	2018-01-03 00:29:37	082567374418739	297	379321864695232
4389973676463558	2018-01-26 13:47:46	295554828848966	400	4389973676463558
4407230633003235	2018-01-27 07:21:08	761335698364860	567	4407230633003235
5403923427969691	2018-01-22 23:46:19	922077754605834	324	5403923427969691
5508842242491554	2018-01-31 14:55:58	634200295989311	585	5508842242491554
6562510549485881	2018-01-17 08:35:27	659982919406634	518	6562510549485881
340028465709212	2018-01-02 03:25:35	009250698176266	233	340028465709212
349143706735646	2018-01-29 22:33:14	343824445342591	298	349143706735646
4126356979547079	2018-01-24 16:09:03	015582765997171	345	4126356979547079
4484950467600170	2018-01-10 08:03:13	570539968421790	462	4484950467600170
4818950814628962	2018-01-31 00:53:15	819006616594636	660	4818950814628962
5464688416792307	2018-01-26 19:03:47	434792568351651	469	5464688416792307
5543219113990484	2018-01-13 18:34:00	501241235491851	494	5543219113990484
5573293264792992	2018-01-31 14:55:57	350307876868039	284	5573293264792992
6011273561157733	2018-02-01 01:27:58	314862932674883	411	6011273561157733
6011985140563103	2018-01-30 02:03:54	393165367933607	350	6011985140563103

only showing top 20 rows

Calculating UCL:

- Calculate the moving average and standard deviation of the last 10 transactions for each card_id for the data present in Hadoop and NoSQL database
- With the fresh dataframe, use member ID once again as common key and join with card_transaction.csv to load postcode, pos_id, status, amount & transaction date fields from history transactions
- open a window frame where we group input dataframe rows on card_id and order by transaction date to get all transactions on card in chronological order

```
window = Window.partitionBy(history['card_id']).orderBy(history['transaction_date'].desc())

history_df = history.select('*', f.rank().over(window).alias('rank')).filter(f.col('rank') <= 10)
```

```
history_df.show()
```

card_id	amount	postcode	pos_id	status	score	transaction_date	rank
340379737226464	1784098	26656	000383013889790	GENUINE	229	2018-01-27 00:19:47	1
340379737226464	3759577	61334	016312401940277	GENUINE	229	2018-01-18 14:26:09	2
340379737226464	4080612	51338	562082278231631	GENUINE	229	2018-01-14 20:54:02	3
340379737226464	4242710	96105	285501971776349	GENUINE	229	2018-01-11 19:09:55	4
340379737226464	9061517	40932	232455833079472	GENUINE	229	2018-01-10 20:20:33	5
340379737226464	102248	40932	232455833079472	GENUINE	229	2018-01-10 15:04:33	6
340379737226464	7445128	50455	915439934619047	GENUINE	229	2018-01-07 23:52:27	7
340379737226464	5706163	50455	915439934619047	GENUINE	229	2018-01-07 22:07:07	8
340379737226464	8090127	18626	359283931604637	GENUINE	229	2017-12-29 13:24:07	9
340379737226464	9282351	41859	808326141065551	GENUINE	229	2017-12-28 19:50:46	10
345406224887566	1135534	53034	146838238062262	GENUINE	349	2017-12-25 04:03:58	1
345406224887566	5190295	88036	821406924682103	GENUINE	349	2017-12-20 04:41:07	2
345406224887566	5970187	28334	024341862357645	GENUINE	349	2017-11-30 05:24:25	3
345406224887566	3854486	48880	172521878612232	GENUINE	349	2017-09-21 00:01:58	4
345406224887566	1242240	14510	536497882467098	GENUINE	349	2017-06-11 16:31:45	5
345406224887566	9222549	68358	875905403447795	GENUINE	349	2017-06-10 21:13:03	6
345406224887566	8726784	64487	617331009748827	GENUINE	349	2017-03-16 03:04:40	7
345406224887566	2415599	99137	751829480922658	GENUINE	349	2017-03-08 12:29:44	8
345406224887566	9671941	65614	607206139883123	GENUINE	349	2017-01-21 08:42:47	9
345406224887566	7454950	18249	368724323320131	GENUINE	349	2016-12-30 04:46:01	10

only showing top 20 rows

Import sql function and then calculate Stddev on amount field
UCL i.e. moving average + 3 * (standard deviation)

```
import pyspark.sql.functions as f
```

```
history_df = history_df.groupBy("card_id").agg(f.round(f.avg('amount'),2).alias('moving_avg'), \
                                              f.round(f.stddev('amount'),2).alias('Std_Dev'))
history_df.show()
```

card_id	moving_avg	Std_Dev
340379737226464	5355453.1	3107063.55
345406224887566	5488456.5	3252527.52
348962542187595	5735629.0	3089916.54
377201318164757	5742377.7	2768545.84
379321864695232	4713319.1	3203114.94
4389973676463558	4923904.7	2306771.9
4407230633003235	4348891.3	3274883.95
5403923427969691	5375495.6	2913510.72
5508842242491554	4570725.9	3229905.04
6562510549485881	5551056.9	2501552.48
340028465709212	6863758.9	3326644.65
349143706735646	5453372.9	3424332.26
4126356979547079	4286400.2	2909676.26
4484950467600170	4550480.5	3171538.48
4818950814628962	2210428.9	958307.87
5464688416792307	4985938.2	2379084.95
5543219113990484	4033586.9	2969107.42
5573293264792992	3929994.0	2589503.93
6011273561157733	4634624.8	2801886.17
6011985140563103	5302878.9	3088988.7

```
history_df = history_df.withColumn('UCL',history_df.moving_avg+3*(history_df.Std_Dev))
history_df.show()
```

card_id	moving_avg	Std_Dev	UCL
340379737226464	5355453.1	3107063.55	1.4676643749999998E7
345406224887566	5488456.5	3252527.52	1.524603906E7
348962542187595	5735629.0	3089916.54	1.5005378620000001E7
377201318164757	5742377.7	2768545.84	1.4048015219999999E7
379321864695232	4713319.1	3203114.94	1.432266392E7
4389973676463558	4923904.7	2306771.9	1.1844220399999999E7
4407230633003235	4348891.3	3274883.95	1.4173543150000002E7
5403923427969691	5375495.6	2913510.72	1.411602776E7
5508842242491554	4570725.9	3229905.04	1.4260441020000001E7
6562510549485881	5551056.9	2501552.48	1.305571434E7
340028465709212	6863758.9	3326644.65	1.684369285E7
349143706735646	5453372.9	3424332.26	1.572636968E7
4126356979547079	4286400.2	2909676.26	1.301542898E7
4484950467600170	4550480.5	3171538.48	1.406509594E7
4818950814628962	2210428.9	958307.87	5085352.51
5464688416792307	4985938.2	2379084.95	1.212319305E7
5543219113990484	4033586.9	2969107.42	1.294090916E7
5573293264792992	3929994.0	2589503.93	1.1698505790000001E7
6011273561157733	4634624.8	2801886.17	1.3040283309999999E7
6011985140563103	5302878.9	3088988.7	1.4569845000000002E7

Join latest dataframe with previous to get all the required data. Final lookup table looks as below:

```
lookup_table.show() #Final data set look as below
```

card_id	transaction_date	amount	postcode	pos_id	status	score	UCL
340379737226464	2018-01-27 00:19:47	1784098	26656	000383013889790	GENUINE	229	1.4676643749999998E7
345406224887566	2017-12-25 04:03:58	1135534	53034	146838238062262	GENUINE	349	1.524603906E7
348962542187595	2018-01-29 17:17:14	7408949	27830	453850044027107	GENUINE	522	1.5005378620000001E7
377201318164757	2017-11-28 16:32:22	4799826	84302	287431794718846	GENUINE	432	1.4048015219999999E7
379321864695232	2018-01-03 00:29:37	5702120	98837	638380208258390	GENUINE	297	1.432266392E7
4389973676463558	2018-01-26 13:47:46	7196505	10985	588476547410852	GENUINE	400	1.1844220399999999E7
4407230633003235	2018-01-27 07:21:08	38579	50167	697070998627535	GENUINE	567	1.4173543150000002E7
5403923427969691	2018-01-22 23:46:19	1576154	17350	734614251977032	GENUINE	324	1.411602776E7
5508842242491554	2018-01-31 14:55:58	2710473	12986	990193545769550	GENUINE	585	1.4260441020000001E7
6562510549485881	2018-01-17 08:35:27	5939348	35440	901627725704672	GENUINE	518	1.305571434E7
340028465709212	2018-01-02 03:25:35	8696557	24658	246987608008994	GENUINE	233	1.684369285E7
349143706735646	2018-01-29 22:33:14	9246599	99101	743905143665678	GENUINE	298	1.572636968E7
4126356979547079	2018-01-24 16:09:03	1770784	14475	698032801419746	GENUINE	345	1.301542898E7
4484950467600170	2018-01-10 08:03:13	2284955	13324	653851258729390	GENUINE	462	1.406509594E7
4818950814628962	2018-01-31 00:53:15	2316346	88081	127695801600255	GENUINE	660	5085352.51
5464688416792307	2018-01-26 19:03:47	4067979	71670	111365575664933	GENUINE	469	1.212319305E7
5543219113990484	2018-01-13 18:34:00	549641	62273	039213658608911	GENUINE	494	1.294090916E7
5573293264792992	2018-01-31 14:55:57	4827477	27012	805073498705051	GENUINE	284	1.1698505790000001E7
6011273561157733	2018-02-01 01:27:58	5272574	45305	063916192266113	GENUINE	411	1.3040283309999999E7
6011985140563103	2018-01-30 02:03:54	1725430	36587	914045782120401	GENUINE	350	1.4569845000000002E7

only showing top 20 rows

Drop duplicates on this DF to remove redundant transactions done of card_id, transaction date, score & post code.

```
lookup_table = lookup_table.dropDuplicates(['card_id', 'transaction_date', 'postcode'])
```

```
lookup_table.count() #1000
```

1000

Task4: Load the data in lookup table

Steps:

1. Used hbase as NoSQL database.
2. Open the connection.
3. Created definitions for opening and closing connections.
4. Created definitions for listing, creating and bulk load data.
5. Bulk loaded data in the lookup table in hbase.
6. Used 'list' command in hbase to see the table created.
7. Used scan "lookup_table" to check on the data.
8. Used count "lookup_table" to check on count of data.

```
import happybase
#create connection
connection = happybase.Connection('localhost', port=9090 ,autoconnect=False)
```

```
def open_connection():
    connection.open()
#close the opened connection
def close_connection():
    connection.close()
#list all tables in Hbase
def list_tables():
    print "fetching all table"
    open_connection()
    tables = connection.tables()
    close_connection()
    print "all tables fetched"
    return tables
```

```
#create the required table
def create_table(name,cf):
    print "creating table " + name
    tables = list_tables()
    if name not in tables:
        open_connection()
        connection.create_table(name, cf)
        close_connection()
        print "table created"
    else:
        print "table already present"
#get the pointer to a table
def get_table(name):
    open_connection()
    table = connection.table(name)
    close_connection()
    return table
```

```
create_table('lookup_table', {'info' : dict(max_versions=5) })
```

```
creating table lookup_table
fetching all table
all tables fetched
table created
```

```
#batch insert data in Lookup table
def batch_insert_data(df,tableName):
    print "starting batch insert of events"
    table = get_table(tableName)
    open_connection()
    rows_count=0

    #Creating a rowkey for better data query. RowKey is the cardId .
    rowKey_dict={}
    with table.batch(batch_size=4) as b:
        for row in df.rdd.collect():
            b.put(bytes(row.card_id) , { 'info:card_id':bytes(row.card_id),
                                         'info:transaction_date':bytes(row.transaction_date),
                                         'info:score':bytes(row.score),
                                         'info:postcode':bytes(row.postcode),
                                         'info:UCL':bytes(row.UCL)})

    print "batch insert done"
    close_connection()
```

```
batch_insert_data(lookup_table,'lookup_table')
```

```
starting batch insert of events
batch insert done
```

```
# create table of card_transactions.csv file.
create_table('card_transactions', {'info' : dict(max_versions=5) })
```

```
creating table card_transactions
fetching all table
all tables fetched
table created
```

```
def batch_insert_csvdata(filename,tableName):
    print "starting batch insert of events"
    file = open(filename, "r")
    table = get_table(tableName)
    open_connection()
    i=0

    for line in file:
        temp = line.strip().split(",")

        #Skip the first row
        if temp[0]!='card_id':

            table.put(bytes(i) , { 'info:card_id':bytes(temp[0]),
                                   'info:member_id':bytes(temp[1]),
                                   'info:amount':bytes(temp[2]),
                                   'info:postcode':bytes(temp[3]),
                                   'info:pos_id':bytes(temp[4]),
                                   'info:transaction_dt':bytes(temp[5]),
                                   'info:status':bytes(temp[6])})

            i=i+1

    file.close()
    print "batch insert done"
    close_connection()
```

```
#Batch insert data of card_transactions.csv file.
batch_insert_csvdata('card_transactions.csv','card_transactions')
```

```
starting batch insert of events
batch insert done
```

Validate the table created and data in Hbase:

- 1) Login to putty as root user
- 2) Start thrift server using below command
/opt/cloudera/parcels/CDH/lib/hbase/bin/hbase-daemon.sh start thrift -p 9090
- 3) Give command hbase shell
- 4) Give command "list"


```
hbase(main):007:0> list
TABLE
card_transactions
lookup_table
2 row(s) in 0.0070 seconds

=> ["card_transactions", "lookup_table"]
hbase(main):008:0> █
```

root@ip-10-0-0-149:~

```
6591175617713393      column=info:transaction_date, timestamp=1634375007372, value=2018-01-31 13:10:37
6592184145413632      column=info:UCL, timestamp=1634375006787, value=13734342.65
6592184145413632      column=info:card_id, timestamp=1634375006787, value=6592184145413632
6592184145413632      column=info:postcode, timestamp=1634375006787, value=53186
6592184145413632      column=info:score, timestamp=1634375006787, value=456
6592184145413632      column=info:transaction_date, timestamp=1634375006787, value=2018-01-28 00:54:30
6594248319343442      column=info:UCL, timestamp=1634375006872, value=15065362.77
6594248319343442      column=info:card_id, timestamp=1634375006872, value=6594248319343442
6594248319343442      column=info:postcode, timestamp=1634375006872, value=24927
6594248319343442      column=info:score, timestamp=1634375006872, value=350
6594248319343442      column=info:transaction_date, timestamp=1634375006872, value=2018-01-31 23:42:38
6595638658736751      column=info:UCL, timestamp=1634375007621, value=14005069.97
6595638658736751      column=info:card_id, timestamp=1634375007621, value=6595638658736751
6595638658736751      column=info:postcode, timestamp=1634375007621, value=68328
6595638658736751      column=info:score, timestamp=1634375007621, value=310
6595638658736751      column=info:transaction_date, timestamp=1634375007621, value=2018-01-30 10:50:34
6595814135833988      column=info:UCL, timestamp=1634375007288, value=14332708.84
6595814135833988      column=info:card_id, timestamp=1634375007288, value=6595814135833988
6595814135833988      column=info:postcode, timestamp=1634375007288, value=22508
6595814135833988      column=info:score, timestamp=1634375007288, value=210
6595814135833988      column=info:transaction_date, timestamp=1634375007288, value=2018-01-30 02:03:54
6595928469079750      column=info:UCL, timestamp=1634375008323, value=11824730.01
6595928469079750      column=info:card_id, timestamp=1634375008323, value=6595928469079750
6595928469079750      column=info:postcode, timestamp=1634375008323, value=98349
6595928469079750      column=info:score, timestamp=1634375008323, value=412
6595928469079750      column=info:transaction_date, timestamp=1634375008323, value=2018-01-24 12:38:22
6597703848279563      column=info:UCL, timestamp=1634375007681, value=15250624.49
6597703848279563      column=info:card_id, timestamp=1634375007681, value=6597703848279563
6597703848279563      column=info:postcode, timestamp=1634375007681, value=95699
6597703848279563      column=info:score, timestamp=1634375007681, value=218
6597703848279563      column=info:transaction_date, timestamp=1634375007681, value=2018-01-27 10:51:49
6598830758632447      column=info:UCL, timestamp=1634375007878, value=12685782.48
6598830758632447      column=info:card_id, timestamp=1634375007878, value=6598830758632447
6598830758632447      column=info:postcode, timestamp=1634375007878, value=19421
6598830758632447      column=info:score, timestamp=1634375007878, value=293
6598830758632447      column=info:transaction_date, timestamp=1634375007878, value=2018-01-30 00:18:34
6599900931314251      column=info:UCL, timestamp=1634375008288, value=12487392.07
6599900931314251      column=info:card_id, timestamp=1634375008288, value=6599900931314251
6599900931314251      column=info:postcode, timestamp=1634375008288, value=97423
6599900931314251      column=info:score, timestamp=1634375008288, value=297
6599900931314251      column=info:transaction_date, timestamp=1634375008288, value=2018-01-31 11:25:16
999 row(s) in 0.4250 seconds
hbase(main):009:0> █
```


Validation as per mentioned:

Count of card_transactions.csv: 53292

```
: transaction = StructType([StructField('card_id', StringType(),False),
                               StructField('member_id', StringType(),False),
                               StructField('amount', IntegerType(),False),
                               StructField('postcode', StringType(),False),
                               StructField('pos_id', StringType(),False),
                               StructField('transaction_dt', StringType(),False),
                               StructField('status', StringType(),False),
                               ])

: tranf = spark.read.csv("hdfs://user/root/cap_project/card_transactions.csv", header = True, schema = transaction)

: tranf.count()

: 53292
```

Count of sqoop jobs :999

```
cardschema = StructType([StructField('card_id', StringType(),False),
                               StructField('member_id', StringType(),False),
                               StructField('member_joining_dt', StringType(),False),
                               StructField('card_purchase_dt', StringType(),False),
                               StructField('country', StringType(),False),
                               StructField('city', StringType(),False),
                               ])

cardf = spark.read.csv("hdfs://user/root/cap_project/card_member", header = False, schema = cardschema)

#Checking number of records loaded from HDFS, Looks they are loaded as expected.
cardf.count()

999
```

```
memberschema = StructType([StructField('member_id', StringType(),False),
                               StructField('score', IntegerType(),False),
                               ])

memf = spark.read.csv("hdfs://user/root/cap_project/member_score", header = False, schema = memberschema)

memf.count()

999
```